




```
import pandas as pd
import numpy as np

df = pd.read_csv('/content/Mall_Customers.csv')
df
```






	CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40
...
195	196	Female	35	120	79
196	197	Female	45	126	28
197	198	Male	32	126	74
198	199	Male	32	137	18
199	200	Male	30	137	83

200 rows × 5 columns

```
df.columns

Index(['CustomerID', 'Genre', 'Age', 'Annual Income (k$)',
      'Spending Score (1-100)'],
      dtype='object')
```

```
df.head(10)
```



	CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40
5	6	Female	22	17	76
6	7	Female	35	18	6
7	8	Female	23	18	94
8	9	Male	64	19	3
9	10	Female	30	19	72

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            200 non-null   int64
1   Genre                 200 non-null   object
2   Age                   200 non-null   int64
3   Annual Income (k$)    200 non-null   int64
4   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
df.describe()
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)	
count	200.000000	200.000000	200.000000	200.000000	
mean	100.500000	38.850000	60.560000	50.200000	
std	57.879185	13.969007	26.264721	25.823522	
min	1.000000	18.000000	15.000000	1.000000	
25%	50.750000	28.750000	41.500000	34.750000	
50%	100.500000	36.000000	61.500000	50.000000	
75%	150.250000	49.000000	78.000000	73.000000	
max	200.000000	70.000000	137.000000	99.000000	

df.isnull().sum()

CustomerID 0
Genre 0
Age 0
Annual Income (k\$) 0
Spending Score (1-100) 0
dtype: int64

df.nunique()

CustomerID 200
Genre 2
Age 51
Annual Income (k\$) 64
Spending Score (1-100) 84
dtype: int64

df.dtypes

CustomerID int64
Genre object
Age int64
Annual Income (k\$) int64
Spending Score (1-100) int64
dtype: object

df == 0

	CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)	
0	False	False	False	False	False	
1	False	False	False	False	False	
2	False	False	False	False	False	
3	False	False	False	False	False	
4	False	False	False	False	False	
...	
195	False	False	False	False	False	
196	False	False	False	False	False	
197	False	False	False	False	False	
198	False	False	False	False	False	
199	False	False	False	False	False	

200 rows × 5 columns

df[df==0]


```
CustomerID      0
Genre           0
Age            0
Annual Income (k$) 0
Spending Score (1-100) 0
dtype: int64
```

(200, 5)

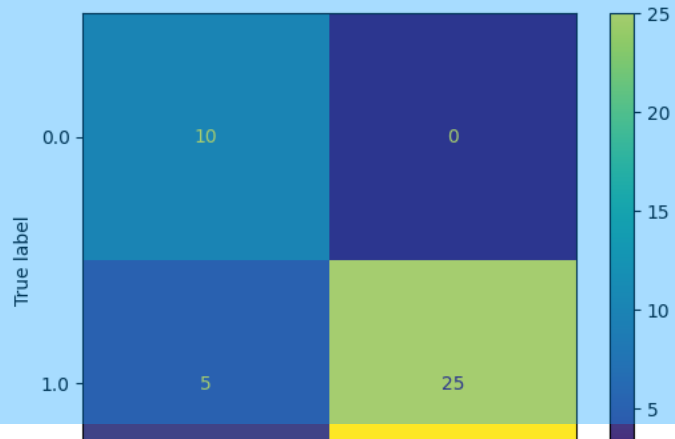
```
x = df[['Annual Income (k$)', 'Age']]
y = df['Spending Score (1-100)']
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=1)
```

 $(160, 2)$ $(40, 2)$ [illegible][illegible]

numpy.ndarray

```
ConfusionMatrixDisplay.from_predictions(actual, predicted)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f3cf2472d40>
```



```
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
```

```

              0.0              1.0
print(classification_report(actual,predicted))
```

	precision	recall	f1-score	support
0.0	0.67	1.00	0.80	10
1.0	1.00	0.83	0.91	30
accuracy			0.88	40
macro avg	0.83	0.92	0.85	40
weighted avg	0.92	0.88	0.88	40

```
accuracy_score(actual,predicted)
```

```
0.875
```